

# **Scientific Computation**

**Spring, 2019**

**Lecture 14**

# Notes

---

- **Extra office hour this week: Thursday, 4-5pm, MLC**

## **Plan for last few weeks:**

- **Lectures 15 and 16: Discrete Fourier transforms and nonlinear time series analysis**
- **Lectures 17 and 18: Numerical differentiation**
- **Lectures 19 and 20: Compiling Python code, parallel computing with Python, suggested topics from class**

# Data analysis

---

- With PCA, the goal was to extract information from a dataset
- Now, we want to think about filling in information *missing* from a dataset
- There are many different approaches, we will look at just one based on low-rank matrix factorization
- Is this actually useful? Very important in computer graphics, vision, machine learning, and recommender systems
- We will focus on the latter example

# Recommender systems

---

- How do Netflix, Amazon, Facebook, etc... decide what to recommend to you?
- They collect:
  - information about what you like (and dislike)
  - information about what everyone else likes
- And then attempt to predict what you will spend time and/or money on
- How is this data organized? – one example is a ratings matrix

# Recommender systems

---

- How is this data organized? – one example is a ratings matrix

	Suits	Sex Education	Friends	Stranger Things	Killing Eve
Don	5	?	1	?	?
Liz	0	4	5	?	?
Kamala	2	?	3	?	5
Beto	1	5	4	5	?

- How do we fill in the missing entries?
- Organizing idea: Need ratings for high-level concepts (e.g. genre) rather than individual movies

- Note: in practice, ratings matrices are much, much larger

# Recommender systems

---

- Let's first think about a simpler case – how do we fill in the missing entries in this matrix:  $D = \begin{bmatrix} 1 & 2 \\ \times & 6 \\ 2 & \times \end{bmatrix}$
- We need to set criteria to assess how well we fill in the data
- Basic idea: Fill in the data without increasing the “information” or introducing “new trends”

# Recommender systems

---

- Let's first think about a simpler case – how do we fill in the missing entries in

this matrix:

$$D = \begin{bmatrix} 1 & 2 \\ \times & 6 \\ 2 & \times \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 2 \\ 3 & 6 \\ 2 & 4 \end{bmatrix}$$

- We need to set criteria to assess how well we fill in the data
- Basic idea: Fill in the data without increasing the “information” or introducing “new trends”
- If we maximized variance like in PCA, we would be doing the opposite and inventing trends
- We could think about minimizing variance, but a simpler closely-related idea is to minimize the *rank*
  - $\text{Rank}(D)$  = dimension of column space of  $D$  = dimension of row space of  $D$
  - In the example above, the rank-1 estimate would be as above

# Rank minimization

---

$$D = \begin{bmatrix} 1 & 2 \\ \times & 6 \\ 2 & \times \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 2 \\ 3 & 6 \\ 2 & 4 \end{bmatrix}$$

- With this approach, we have estimated “ratings” without introducing new information or trends
- Basic idea of matrix factorization-based recommender systems:
  - Complete the ratings matrix so that:
    - Changes to existing entries are “small”
    - The rank of the resulting matrix is minimized



# Rank minimization

---

$$D = \begin{bmatrix} 1 & 2 \\ \times & 6 \\ 2 & \times \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 2 \\ 3 & 6 \\ 2 & 4 \end{bmatrix}$$

- With this approach, we have estimated “ratings” without introducing new information or trends
- Basic idea of matrix factorization-based recommender systems:
  - Complete the ratings matrix so that:
    - Changes to existing entries are “small”
    - The rank of the resulting matrix is minimized
  - We will look at an approach based on the SVD
    - A SVD is a sum of rank-1 matrices
    - $\text{Rank}(A)$  = number of non-zero singular values of A

# Rank minimization

---

- This is an optimization problem, and we first need to define a *cost function*
- Motivating example:  
  
Find “completed” ratings matrix,  $A$ , such that  $\text{rank}(A)$  is minimized and  $A = R$  for all valid entries in  $R$
- The rank is discrete, and it is easier to solve continuous optimization problems with differentiable cost functions (NP hard to minimize rank while leaving existing ratings unchanged)
- Instead of minimizing rank, we can maximize the sum of the singular values of  $A$  (the rank is the number of non-zero singular values):

Find  $A$  such that  $|A|_*$  is minimized and  $A_{i,j} = R_{i,j}$  for  $(i, j) \in \Omega$ .

$\Omega$  is the set of indices for which ratings have been specified in  $R$ , and  $|A|_*$  is the sum of the singular values of  $A$ .

# Rank minimization

---

- For numerical calculations, the missing ratings are typically replaced with a number, say, -1000 or something that cannot be mistaken for an actual rating
- We then need an auxiliary matrix to help “remove” these numbers:

$$E_{i,j} = \begin{cases} 0, & (i, j) \in \Omega \\ R_{i,j} - A_{i,j} & (i, j) \notin \Omega \end{cases}$$

- So, our optimization problem requires:  $R - A - E = 0$
- How do we enforce this constraint? With a matrix of *Lagrange multipliers*,  $Y$ , which must be found as part of the problem
- Our provisional cost function (Lagrangian) is:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij})$$

# Rank minimization

---

- Our provisional cost function (Lagrangian) is:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij})$$

- We add one more *regularizing* term:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- This reduces the likelihood of large ratings being placed in A
- And introduces a new parameter,  $\mu$ , that must be determined
- But it also allows us to leverage a powerful optimization result (shown later)

# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A, E, Y, \mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- Next step: construct method to solve the optimization problem.

# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A$ ,  $E$ ,  $Y$ ,  $\mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- Next step: construct method to solve the optimization problem.
- We could use any standard gradient-based optimization method (e.g. a optimizer in `scipy.minimize`)
  - This requires providing the derivatives of the Lagrangian with respect to the entries of  $A$ ,  $E$ ,  $Y$ , and  $\mu$
- We will take a different approach here. We use an iterative approach and update the four parameters sequentially. This works particularly well for large ratings matrices with low rank

# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A, E, Y, \mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu \|R - A - E\|_F$$

- Step 0: Provide initial guesses for  $A, E, Y, \mu$
- Given estimates for these parameters at iteration  $k$ , there is a four-stage process to obtain improved estimates at iteration  $k+1$
- Stage 1:  $A_{k+1} = \arg \min_A \mathcal{L}(A, E_k, Y_k, \mu_k)$
- “Find  $A$  such that the Lagrangian is minimized with  $E, Y, \mu$  held fixed to their values at iteration  $k$ ”
- A recent result from optimization theory – this problem is solved with:

$$\text{svd}(R - E_k + \frac{1}{\mu_k} Y_k) = U \Sigma V^T$$

$$A_{k+1} = U \tilde{\Sigma}_{\mu_k} V^T$$

# Rank minimization

---

- **Stage 1:**  $A_{k+1} = \arg \min_A \mathcal{L}(A, E_k, Y_k, \mu_k)$
- “Find A such that the Lagrangian is minimized with , E, Y,  $\mu$  held fixed to their values at iteration k”
- A recent result from optimization theory – this problem is solved with:

$$\text{svd}(R - E_k + \frac{1}{\mu_k} Y_k) = U \Sigma V^T$$

$$A_{k+1} = U \tilde{\Sigma}_{\mu_k} V^T$$

- Here, we are iteratively “reducing” the singular values of A (and reducing its rank) with:  $\tilde{\sigma}^{(i)} = \max(\sigma^{(i)} - \frac{1}{\mu_k}, 0)$
- So singular values of  $(R - E_k + \frac{1}{\mu_k} Y_k)$  which are smaller than  $1/\mu$  are set to zero, and all others are reduced by  $1/\mu$



# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A, E, Y, \mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- **Step 0:** Provide initial guesses for  $A, E, Y, \mu$
- Given estimates for these parameters at iteration  $k$ , there is a four-stage process to obtain improved estimates at iteration  $k+1$
- **Stage 2:**  $E_{k+1} = \arg \min_E \mathcal{L}(A_{k+1}, E, Y_k, \mu_k)$
- Differentiating the Lagrangian with respect to  $E$  and setting the result to zero:

$$E_{k+1} = \frac{1}{2\mu_k} Y_k + R - A_{k+1}, \quad (i, j) \notin \Omega$$

# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A$ ,  $E$ ,  $Y$ ,  $\mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- **Step 0:** Provide initial guesses for  $A$ ,  $E$ ,  $Y$ ,  $\mu$
- Given estimates for these parameters at iteration  $k$ , there is a four-stage process to obtain improved estimates at iteration  $k+1$
- **Stage 3:**  $Y_{k+1} = Y_k + \mu_k (D - A_{k+1} - E_{k+1})$ .
- This is a *relaxation* of  $Y$  based on how well constraint is satisfied and is a companion result to what we used for stage 1

# Rank minimization

---

- We can now state our “full” optimization problem:  
Find  $A$ ,  $E$ ,  $Y$ ,  $\mu$ , such that the Lagrangian (below) is minimized:

$$\mathcal{L} = |A|_* + \sum_{i=1}^M \sum_{j=1}^N Y_{ij} (R_{ij} - A_{ij} - E_{ij}) + \mu |R - A - E|_F$$

- **Step 0:** Provide initial guesses for  $A$ ,  $E$ ,  $Y$ ,  $\mu$
- Given estimates for these parameters at iteration  $k$ , there is a four-stage process to obtain improved estimates at iteration  $k+1$
- **Stage 4:**  $\mu_{k+1} = \rho \mu_k$ ,  $\rho > 1$
- At larger  $\mu$ , our approximation is closer to the true minimum of the Lagrangian, so we gradually increase  $\mu$  each iteration. Increase more gradually if there are more missing entries

# Rank minimization

---

- See *rec.py* for Python implementation of this method
- Applying it to our small example:

	Suits	Sex Education	Friends	Stranger Things	Killing Eve
Don	5	1	1	1	2
Liz	0	4	5	4	3
Kamala	2	4	3	4	5
Beto	1	5	4	5	5

# Recommender systems

---

- There are a broad range of matrix factorization approaches
  - For example, estimate  $r$  column vectors and  $r$  row vectors that minimize an error based on Frobenius norm
- Different optimization methods can then be applied for each of these approaches
  - Stochastic gradient descent is popular for very large ratings matrices
- Underlying idea, as in PCA, and linear amplification problems: singular values and corresponding eigenvectors provide essential information about key trends “hidden” in data